

# Linux Operating System

Quang Hoang - 2024

Viettel Digital Talent 2024 - Software & Data Engineering

# Lesson Outline

1. **Linux Kernel deep dive.**
  - **Introduction**
  - **Process**
  - **Kernel synchronization**
  - **Virtual Filesystem**
2. **Linux command lines.**

# Why Linux?

- Performance
  - Fast and Stable
  - No requirement for latest hardware
- Security
- Customizability
- It is FREE
  - GNU General Public License (GPL)
    - free to **download** the source code and make any **modifications**
    - if you distribute your changes, it should be GPL license as well.

# What is Linux?

- Before Linux
  - MS DoS was the dominated OS for PC: single user, closed source.
  - Apple MacOS was better but expensive.
  - UNIX was much better (multitasking and multi-user) but much expensive.
  - People were looking for UNIX-like OS which is cheaper.

# What is Linux?

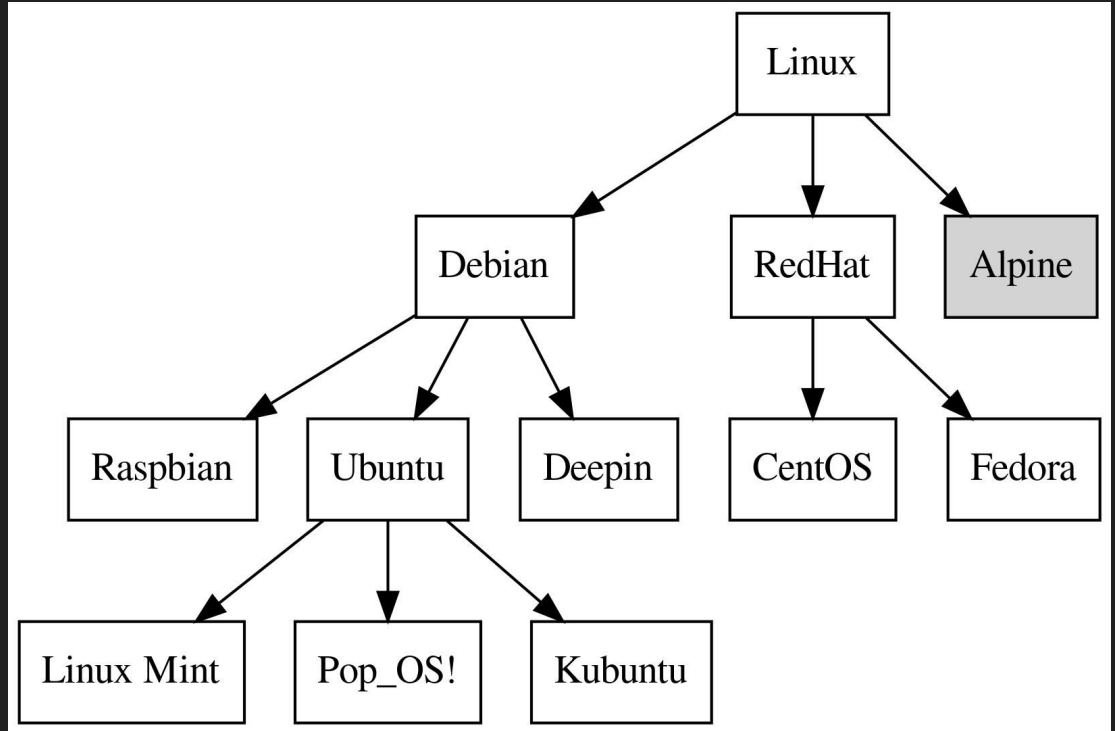
- Before Linux
  - MS DoS was the dominated OS for PC: closed source.
  - Apple MAC was better but expensive.
  - UNIX was much better but much expensive.
  - People were looking for UNIX-like OS which is cheaper.
- Beginning of Linux
  - 1991: Linus Torvalds developed Linux v0.0.1

# What is Linux?

- Before Linux
  - MS DoS was the dominated OS for PC: closed source.
  - Apple MAC was better but expensive.
  - UNIX was much better but much expensive.
  - People were looking for UNIX-like OS which is cheaper.
- Beginning of Linux
  - 1991: Linus Torvalds developed Linux v0.0.1
- Linux today
  - Power smartphone, smartwatch, PC, Supercomputers
    - 71.9% of smartphone using Android as of 2022
    - 96.55% of web server run Linux as of 2015

# Linux “Distro”

- short for distribution
- modification version of Linux kernel.
- [full list](#)



# Operating System and Kernel

OS is a program:

- manage and protect hardware (HW) resource.
- provide HW resource to application via well-defined APIs.




# Operating System and Kernel

OS is a program:

- manage and protect hardware (HW) resource.
- provide HW resource to application via well-defined APIs.

## Kernel

- core of OS.



# The Process

## Process

- A program in the midst of execution.
- Act like a mini virtual machine:
  - CPU: time slicing of processor.
  - Memory: address space.

## Thread

- A thread is an object of activity within the process. Smallest sequence of programmed instructions that can be managed independently by a scheduler
- One process can have multiple threads.

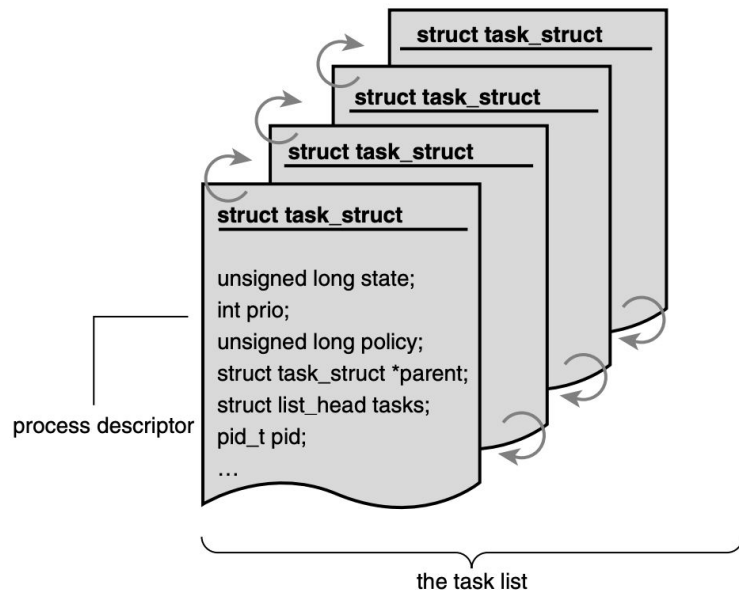
## In Linux

- there is no differentiation between a **process** and a **thread**. Thread is just a special kind of process.

# The Process

Process and Thread in Linux are implemented using the same data structure

A doubly-linked list of **task\_struct**



# Process life-cycle

- Process creation

Init process: the first process created at system boot up. pid = 0.

2 steps:

- fork(): create a child process that is a copy of current process
- exec(): load new executable into the address space and begins executing

# Process life-cycle

- Process creation

```
main() {  
    printf("Hello\n");  
    fork();  
    printf("world\n");  
}
```

# Process life-cycle

- Process creation

What is the output?

Hello

world

world

```
main() {  
    printf("Hello\n");  
    fork();  
    printf("world\n");  
}
```

# Process life-cycle

- Process creation

How many process were created?

```
int main()
{
    int i;
    for(i=0;i<4;i++)
        fork();
    return 0;
}
```

# Process life-cycle

- Process creation

fork();

fork();

fork();

fork();

```
int main()
{
    int i;
    for(i=0;i<4;i++)
        fork();
    return 0;
}
```



# Process life-cycle

- Process creation

p0

fork(); p1

fork(); p2, p3

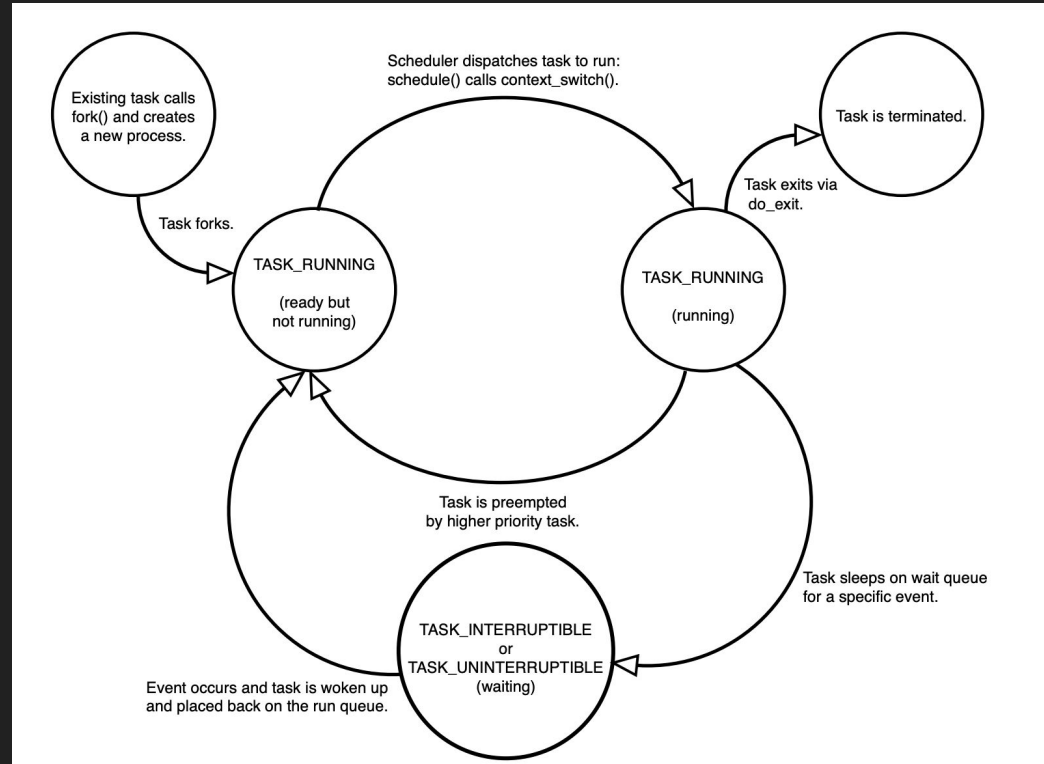
fork(); p4, p5, p6, p7

fork(); p8, p9, p10, p11, p12, p13, p14, p15

```
int main()
{
    int i;
    for(i=0; i<4; i++)
        fork();
    return 0;
}
```

# Process life-cycle

- runnable
- running
- sleeping (waiting)
- stopped



# Process life-cycle

Linux is preemptive

- No matter what, scheduler will force a process to stop if there is a higher priority task pending.
- running => runnable => running

# Process life-cycle

## Process termination

- when a process terminates, kernel releases resource owned by process and notifies its parent.
- a process terminates when:

# Process life-cycle

## Process termination

- when a process terminates, kernel releases resource owned by process and notifies its parent.
- a process terminates when:
  - the process call `exit()` system call. (self-termination)
  - unhandled exception.

# Process life-cycle

## Process termination

- What if parent process is terminated before its children?

# Process life-cycle

## Process termination

- What if parent process is terminated before its children?
  - child processes are called **orphan** processes. (A common misconception is that child processes are by default killed when their parent exits)
  - reparent to init process.

# Process scheduling

- Process scheduler decides which process runs, when, and for how long.
- 2 types of process:
  - **I/O bound: process spend much time submitting and waiting on I/O requests (e.g. access hard disk, wait for network, keyboard input ...)**
    - GUI applications



# Process scheduling

- Process scheduler decides which process runs, when, and for how long.
- 2 types of process:
  - I/O bound: process spend much time submitting and waiting on I/O requests (e.g. access hard disk, wait for network, keyboard input ...)
    - GUI applications
  - **CPU bound: process spend much time executing code.**
    - executing an infinite loop
    - math calculation application: MATLAB, ssh-keygen

Linux favors I/O-bound processes over CPU-bound processes.

# Process scheduling

Goal of process scheduler:

1. **low latency**: fast process response time.
2. **high throughput**: maximum system utilization.

Timeslice: a number presents how long a process can run until it is preempted.

- long timeslice => high latency
- short timeslice => waste time on context switching.

# Completely Fair Scheduler

Idea: device CPU time equally among processes:

N processes, each should get  $(100/N)\%$  of CPU time.

Process	Time
A	8ms
B	4ms
C	16ms
D	4ms



# Completely Fair Scheduler (CFS)

How it works:

- CFS will run each process for some amount of time, selecting next the process that has run the least.
- The runtime of each process is proportional to its priority.
- When CFS is deciding what process to run next, it picks the process with the smallest runtime.
  - use a Red-Black Tree to store list of runnable processes.
  - Pick smallest node in  $O(1)$
  - Insert in  $O(\log n)$

# Kernel Synchronization

## Race condition

What is the output?

```
define foo_g = 0
```

```
function Foo:
```

```
    define foo = foo_g
```

```
    foo = foo + 1
```

```
    foo_g = foo
```

```
function main:
```

```
    create a thread to run Foo
```

```
    create a thread to run Foo
```

```
    wait for all the thread finish
```

```
    if foo_g equal to 2
```

```
        print "expected"
```

```
    else
```

```
        print "unexpected"
```

# Race condition

The situation when 2 or more threads update the share resource at the same time lead to the uncertain state of the shared resource.

The code that access and update the shared data is called **Critical Section**.

# Cause of Race condition

1. Kernel preemption: one process can preempt another
2. Symmetric multiprocessing: 2 or more processors can execute kernel code at the same time
3. Interrupts: an interrupt can occur any time, interrupting the currently running code.

...

# Spin Lock

- The most common lock in Linux Kernel.

## Thread 1

try to lock data using spin lock

successful: acquired lock

access and update data

unlock

...

## Thread 2

try to lock data using spin lock

failed: waiting...

waiting...

waiting...

successful: acquired lock

access and update data



# Spin Lock

Thread will spin (busy loop) while waiting for the lock to be released.

- we should not hold the spin lock too long to avoid wasting CPU.
- you cannot sleep while holding a spin lock

# Spin Lock

- you cannot sleep while holding a spin lock
- why?

thread A acquires spin lock L, then sleeps.

Thread B tries to acquire L, but because L was acquired by A, B has to wait.

If the computer only has 1 core (or multicores but A and B run on the same core), we have a deadlock:

- B waits for A to release L
- A is sleeping and cannot wake up because B is using CPU for spinning.

# Semaphore

- Unlike Spin lock, semaphores in Linux are sleeping locks.

Thread 1

try to lock data using semaphore

**successed**: acquired lock

access and update data

unlock

...

Thread 2

try to lock data using semaphore

**failed**: **sleep**

**sleeping**...

**sleeping**...

**successed**: acquired lock

access and update data

# Semaphore

- Suitable for locks that are held for a long time.
- Not optimal for short period lock because of the overhead of sleeping, wait queue.

N resources, M threads.

# Counting semaphore

Semaphore allows an arbitrary number of simultaneous lock holder.

2 functions:

- `down()`: acquire the semaphore and decrease count by 1
- `up()`: release the semaphore and increase count by 1.

When count == 0, thread goes to sleep and waits for being woken up.

```
S = new Semaphore(count)
```

```
t1.down()
```

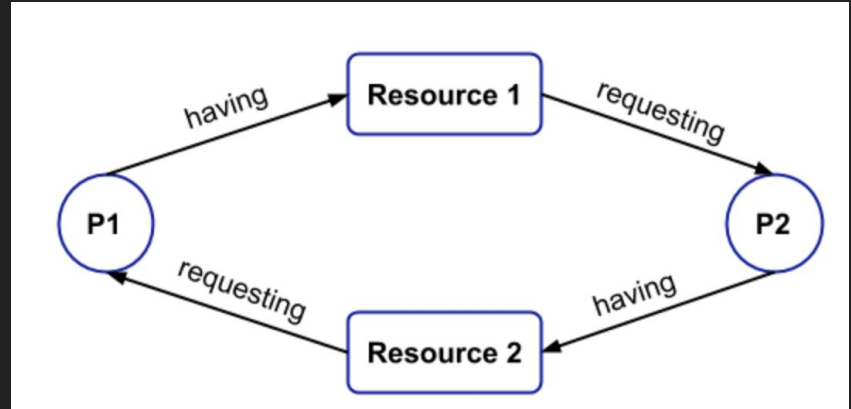

```
t1.up()
```

# Mutex

- mutual exclusion lock
- a sleep lock
- work like a binary semaphore.
- the lock can be acquired and released by the same thread at a time.

# Deadlock

- A deadlock is a situation where a set of processes are blocked because each process is holding a resource and waiting for another resource acquired by some other process.



# Deadlock

- Relock (AA)

P0

```
spin_lock (&A) ;
```

```
...
```

```
spin_lock (&A) ;
```



# Deadlock

- ABBA

P0

```
spin_lock (&A) ;
```

```
...
```

```
spin_lock (&B) ;
```

P1

```
spin_lock (&B) ;
```

```
...
```

```
spin_lock (&A) ;
```

# Deadlock

- ABBCCA deadlocks or more

P0

```
spin_lock (&A) ;
```

```
...
```

```
spin_lock (&B) ;
```

P1

```
spin_lock (&B) ;
```

```
...
```

```
spin_lock (&C) ;
```

P2

```
spin_lock (&C) ;
```

```
...
```

```
spin_lock (&A) ;
```

# Deadlock


## Rules to prevent deadlock

- Implement lock ordering. Nested locks must always be obtained in the same order.
- Do not double acquire the same lock.
- Design for simplicity. Complexity in your locking scheme invites deadlocks.

# Virtual File System (VFS)

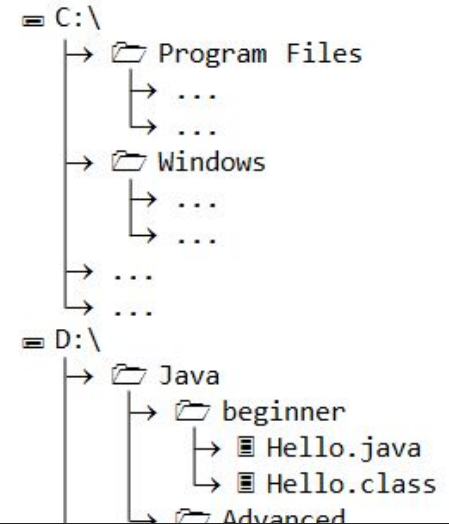

VFS is part of Linux Kernel that implements the file and file-system-related interfaces.

This allows common system calls to work on different type of hardware.



# VFS

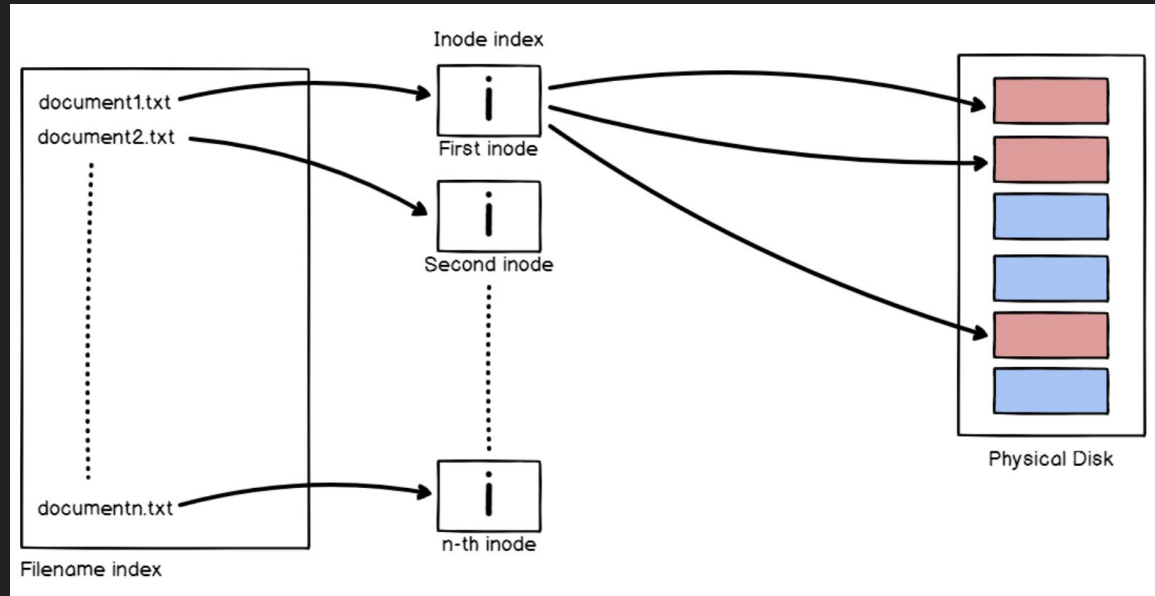
- Single tree
  - unlike Windows
- Mounted filesystem appear as a node in the tree.
- Everything is a file.



# VFS objects

- superblock: represent a specific mounted filesystem
- dentry: represent a directory entry, which is a single component of a path
- inode: represent a specific file, contain metadata (size, owner, location, creation time ...)
  - create
  - link, unlink
  - symlink
- file: represents an open file as associated with a process
  - read
  - write
  - lock
  - open...

# Inode




- to show inode number: `ls -i <file-name>`
- to view data stored in an inode: `stat <file-name>`

# Soft Link, Hard Link

- In Linux, multiple files can point to the same Inode using **Link**.
- Link:
  - a pointer pointing to a file.
  - 2 types:
    - soft link (symbolic link)
    - hard link




# Soft Link (Symbolic Link)



Soft links are files pointing to other file on the file system (like a shortcut in Windows)

- the original file and its soft link share the same content: modify one affects other.
- the original file inode is pointing directly to the file content, soft link inode is pointing to a block containing the path to the original file
- delete or rename original file, the soft link is corrupted.

# Hard Link



Hard link are instances of the file under a different name on the filesystem.

- share the inode of the original file.
- delete/rename original file doesn't affect the hard link.
- hard link vs copy file?

# Prerequisite for next lesson

1. Install [docker](#)
2. Download [this](#) folder
3. Run Ubuntu by docker, access bash terminal and install vim

```
docker run -it ubuntu bash
apt-get update
apt install vim
```

4. Find your ubuntu docker container id

```
docker ps
```

5. Open a new terminal tab, copy VT folder to docker container

```
docker cp VT <id-from-step-4>:/VT
```

6. Verify by running `ls -l VT` in docker terminal.

n processes, sum of memory n processes consume > RAM.

Virtual Memory of OS.

Hard Disk

swap memory

cache invalidation: LRU (Least recently used)